

Benign vs Malignant Findings

Summary

This note outlines the clinical differences between **benign** and **malignant** findings, specifically in breast and lung evaluations.

Benign Findings

- **Growth Pattern:** Non-invasive, localized
- **Borders:** Well-defined, smooth margins
- **Growth Rate:** Slow-growing or stable over time
- **Symptoms:** Often asymptomatic
- **Histology:** Normal cellular architecture, no atypia
- **Examples:**
 - **Breast:** Fibroadenoma, cyst, fibrocystic changes
 - **Lung:** Granuloma, hamartoma, post-inflammatory scar

Impression (Example):

No suspicious mass, distortion, or abnormal calcifications. Findings consistent with benign etiology. Routine follow-up recommended.

Malignant Findings

- **Growth Pattern:** Invasive, potential to spread (metastasis)
- **Borders:** Irregular, spiculated or ill-defined
- **Growth Rate:** Rapid progression
- **Symptoms:** May include pain, weight loss, cough, bleeding

- **Histology:** Atypical cells, mitotic activity, abnormal nuclei
- **Examples:**
 - **Breast:** Invasive ductal carcinoma, lobular carcinoma
 - **Lung:** Adenocarcinoma, squamous cell carcinoma, small cell carcinoma

**Impression (Example):**

Spiculated mass in the upper outer quadrant with associated lymphadenopathy. Findings suspicious for malignancy.

Biopsy recommended. 🧑🏻‍⚕️ Breast Lump Classification: Benign vs Malignant

**Summary**

This project explores the classification of breast tumors into benign (non-cancerous) and malignant (cancerous) types using machine learning models. The goal is to identify which algorithm best predicts tumor type based on various diagnostic features.



Importing the libraries

```
In [27]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import accuracy_score, precision_score, roc_auc_score, ConfusionMatrixDisplay, recall_score, f
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import VotingClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
```

READING IN THE DATASET

```
In [16]: df = pd.read_csv(r'C:\Users\USER\Desktop\breasrcancer_datazset.csv')
df
```

Out[16]:

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	
3	84348301	M	11.42	20.38	77.58	386.1	0.14250	0.28390	
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	
...	
564	926424	M	21.56	22.39	142.00	1479.0	0.11100	0.11590	
565	926682	M	20.13	28.25	131.20	1261.0	0.09780	0.10340	
566	926954	M	16.60	28.08	108.30	858.1	0.08455	0.10230	
567	927241	M	20.60	29.33	140.10	1265.0	0.11780	0.27700	
568	92751	B	7.76	24.54	47.92	181.0	0.05263	0.04362	

569 rows × 33 columns

In [17]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 33 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     569 non-null    int64
1   diagnosis                             569 non-null    object
2   radius_mean                           569 non-null    float64
3   texture_mean                           569 non-null    float64
4   perimeter_mean                         569 non-null    float64
5   area_mean                             569 non-null    float64
6   smoothness_mean                       569 non-null    float64
7   compactness_mean                      569 non-null    float64
8   concavity_mean                        569 non-null    float64
9   concave points_mean                   569 non-null    float64
10  symmetry_mean                         569 non-null    float64
11  fractal_dimension_mean                 569 non-null    float64
12  radius_se                              569 non-null    float64
13  texture_se                             569 non-null    float64
14  perimeter_se                           569 non-null    float64
15  area_se                                569 non-null    float64
16  smoothness_se                          569 non-null    float64
17  compactness_se                         569 non-null    float64
18  concavity_se                           569 non-null    float64
19  concave points_se                      569 non-null    float64
20  symmetry_se                            569 non-null    float64
21  fractal_dimension_se                   569 non-null    float64
22  radius_worst                           569 non-null    float64
23  texture_worst                           569 non-null    float64
24  perimeter_worst                        569 non-null    float64
25  area_worst                             569 non-null    float64
26  smoothness_worst                       569 non-null    float64
27  compactness_worst                      569 non-null    float64
28  concavity_worst                        569 non-null    float64
29  concave points_worst                   569 non-null    float64
30  symmetry_worst                         569 non-null    float64
31  fractal_dimension_worst                 569 non-null    float64
32  Unnamed: 32                             0 non-null      float64
dtypes: float64(31), int64(1), object(1)
memory usage: 146.8+ KB
```

SHAPE OF DATASET

```
In [18]: df.shape
```

```
Out[18]: (569, 33)
```

DISTRIBUTION OF TARGET COLUMN

```
In [19]: df['diagnosis'].value_counts(normalize=True)
```

```
Out[19]: diagnosis
B      0.627417
M      0.372583
Name: proportion, dtype: float64
```

DROPPING OF NULL COLUMN

```
In [20]: df.drop(columns='Unnamed: 32', axis=1, inplace=True)
df.isna().sum()
```

```
Out[20]: id          0
         diagnosis   0
         radius_mean 0
         texture_mean 0
         perimeter_mean 0
         area_mean    0
         smoothness_mean 0
         compactness_mean 0
         concavity_mean 0
         concave points_mean 0
         symmetry_mean 0
         fractal_dimension_mean 0
         radius_se     0
         texture_se     0
         perimeter_se   0
         area_se        0
         smoothness_se  0
         compactness_se 0
         concavity_se   0
         concave points_se 0
         symmetry_se    0
         fractal_dimension_se 0
         radius_worst   0
         texture_worst  0
         perimeter_worst 0
         area_worst     0
         smoothness_worst 0
         compactness_worst 0
         concavity_worst 0
         concave points_worst 0
         symmetry_worst 0
         fractal_dimension_worst 0
         dtype: int64
```


MAPPING TARGET COLUMN TO INT DATA TYPE

```
In [21]: df['diagnosis'] = df['diagnosis'].map({'B':0, 'M':1})
```

CORRELATION MATRIX

```
In [22]: dcorr = df.corr()
plt.figure(figsize=(18, 8))
sns.heatmap(dcorr, annot=True, cmap='coolwarm', fmt='.2f', linewidths=0.8)
```

Out[22]: <Axes: >



SPLITTING THE DATASET

```
In [ ]: X = df.drop(columns=['id', 'diagnosis'], axis=1) # predictor variables
y = df['diagnosis'] # target variables

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, stratify=y, random_state=42) # splitting
```

```
print(X.shape)
print(y.shape)
X_train.head()
```

```
(569, 30)
```

```
(569,)
```

```
Out[ ]:
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	cpoint
78	20.18	23.97	143.70	1245.0	0.12860	0.34540	0.37540	
330	16.03	15.51	105.80	793.2	0.09491	0.13710	0.12040	
378	13.66	15.15	88.27	580.6	0.08268	0.07548	0.04249	
213	17.42	25.56	114.50	948.0	0.10060	0.11460	0.16820	
89	14.64	15.24	95.77	651.9	0.11320	0.13390	0.09966	

5 rows × 30 columns

DUE TO SOMEW OF THE MODEL ARE SENSITIVE I DECIDED TO SCALE THE DATASET USING STANDARD SCALER

```
In [24]: scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

IMPLEMENTED 4 DIFFERENT MODELS TO CHECK WHICH MODEL IS BEST

```
In [ ]: lr = LogisticRegression(random_state=42)
knn = KNeighborsClassifier()
rf = RandomForestClassifier(random_state=42)
dt = DecisionTreeClassifier(random_state=42)

# classification models
classifier = [('K nearest neighbor ', knn),
              ('logistic Regression', lr),
              ('Random Forest', rf),
              ('Decision Tree', dt)]
```



```

for class_name, classes in classifier :

    # cross validation score
    scores = cross_val_score(classes, X_scaled, y_train, cv=5, scoring='accuracy')
    print("Cross-validated scores:", scores)
    print("Mean Accuracy: {:.2f}".format(scores.mean()))

    # fitting all the classes
    classes.fit(X_scaled, y_train)

    # predicting the test set
    y_pred_tree = classes.predict(X_test_scaled)

    print('{:s} : {:.2f}'.format(class_name, accuracy_score(y_test, y_pred_tree)))

```

```

Cross-validated scores: [0.9875      0.9875      0.975      0.98734177 0.91139241]
Mean Accuracy: 0.97
K nearest neighbor : 0.96
Cross-validated scores: [0.9625      1.          0.9875      0.97468354 0.93670886]
Mean Accuracy: 0.97
logistic Regression : 0.97
Cross-validated scores: [0.975      0.9875      0.95      0.96202532 0.88607595]
Mean Accuracy: 0.95
Random Forest : 0.96
Cross-validated scores: [0.9625      0.9625      0.925      0.86075949 0.87341772]
Mean Accuracy: 0.92
Decision Tree : 0.90

```

KNEIGHBORS HAPPENS TO HAVE THE BEST ACCURACY AND CROSS VAL SCORE

```

In [39]: acc_train = knn.score(X_scaled, y_train) # training set accuracy score
         acc_test = knn.score(X_test_scaled, y_test) # testing set accuracy score

         print(f"Knearest Neighbor - Train Accuracy: {acc_train:.2f}")
         print(f"Knearest Neighbor - Test Accuracy: {acc_test:.2f}")

         # predicting the test set
         y_pred = knn.predict(X_test_scaled)

         # cross validation score
         scores = cross_val_score(classes, X_scaled, y_train, cv=10, scoring='accuracy')

```

```
print("Cross-validated scores:", scores)
print('-----')

# Evaluation metrics
accuracy_tree = accuracy_score(y_test, y_pred)

roc_auc_tree = roc_auc_score(y_test, y_pred)

precision_tree = precision_score(y_test, y_pred)

recall_tree = recall_score(y_test, y_pred)

f1_tree = f1_score(y_test, y_pred)

print(f"Knearest Neighbor - Accuracy: {accuracy_tree:.2f}")
print('-----')
print(f"Knearest Neighbor - ROC AUC: {roc_auc_tree:.2f}")
print('-----')
print(f"Knearest Neighbor - Precision: {precision_tree:.2f}")
print('-----')
print(f"Knearest Neighbor - Recall: {recall_tree:.2f}")
print('-----')
print(f"Knearest Neighbor - F1 Score: {f1_tree:.2f}")
print('-----')

print("\nClassification Report:")
print(classification_report(y_test, y_pred))
ConfusionMatrixDisplay.from_estimator(knn, X_test_scaled, y_test, cmap='Blues')
```

```
Knearest Neighbor - Train Accuracy: 0.97
Knearest Neighbor - Test Accuracy: 0.96
Cross-validated scores: [0.95      0.975      0.95      0.975      0.875      0.925
 0.9      0.925      0.84615385 0.87179487]
```

```
-----
Knearest Neighbor - Accuracy: 0.96
-----
```

```
Knearest Neighbor - ROC AUC: 0.95
-----
```

```
Knearest Neighbor - Precision: 1.00
-----
```


```
Knearest Neighbor - Recall: 0.91
-----
```

```
Knearest Neighbor - F1 Score: 0.95
-----
```

Classification Report:

	precision	recall	f1-score	support
0	0.95	1.00	0.97	107
1	1.00	0.91	0.95	64
accuracy			0.96	171
macro avg	0.97	0.95	0.96	171
weighted avg	0.97	0.96	0.96	171

Out[39]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x27795b91310>



SAVING THE MODEL

```
In [43]: import pickle
with open('breast_model.pkl', 'wb') as file:
    pickle.dump(knn, file)
```